

AADHISH S

Generative AI Engineer | LLM Application Developer

Chennai, India | 9544012187 | aadhishs007@gmail.com

linkedin.com/in/aadhish-s | github.com/sAadhish

PROFESSIONAL SUMMARY

Generative AI Engineer with hands-on experience building production-grade LLM applications, RAG pipelines, and semantic search systems. Built JobLens — an AI-powered job intelligence platform using Large Language Models (LLMs), vector databases, embeddings, retrieval-augmented generation (RAG), and LangChain. Skilled in prompt engineering, structured output extraction, and end-to-end AI system design. Prior production experience delivering enterprise data solutions for global stakeholders with strong ability to own projects end-to-end from architecture to deployment.

TECHNICAL SKILLS

Generative AI & LLMs: Large Language Models (LLMs), Prompt Engineering, Retrieval-Augmented Generation (RAG), LLM Chains, AI Agents, Tool Use, Structured Output, Temperature Tuning, Groq API, Llama 3.3 70B

Embeddings & Search: Sentence Transformers (all-MiniLM-L6-v2), Semantic Search, Text Embeddings, Cosine Similarity, Vector Similarity

Frameworks & Libraries: LangChain, FastAPI, Sentence Transformers, python-dotenv

Vector Databases: ChromaDB, HNSW Indexing, ANN (Approximate Nearest Neighbor) Search, Metadata Filtering, Persistent Storage

Programming & Tools: Python, SQL, JSON, Regex, Git, GitHub, REST APIs, Modular Architecture

Data & DevOps: Snowflake, Docker, MySQL

Practices: Error Handling, Graceful Degradation, Logging, Data Validation, UAT/SIT, Agile (JIRA), System Design, Production Debugging

AI PROJECTS

JobLens — AI-Powered Job Market Intelligence Platform

github.com/sAadhish/joblens

Python, Groq API, Llama 3.3 70B, LangChain, ChromaDB, Sentence Transformers, RAG, FastAPI, Docker

- Designed and built a production-grade Generative AI system that performs semantic job matching, structured JD analysis, and AI-powered skill gap assessment using Large Language Models (Groq / Llama 3.3 70B) and retrieval-augmented generation (RAG)
- Implemented a complete RAG pipeline — ingesting job descriptions, chunking text, generating embeddings via Sentence Transformers (all-MiniLM-L6-v2), storing vectors in ChromaDB, and retrieving contextually relevant documents for LLM-powered analysis
- Built a semantic search engine using vector embeddings with cosine similarity ranking, achieving accurate candidate-to-job matching with configurable confidence thresholds and metadata filtering (role, location, experience, remote preference)
- Engineered robust prompt engineering patterns — system/user prompt separation, role-based prompting, structured JSON output with schema enforcement, response sanitization (regex-based markdown stripping),

and temperature=0 for deterministic structured outputs

- Designed modular Python architecture following single-responsibility principle: analyzer (LLM extraction), store (embedding + vector storage), matcher (semantic search + filtering), gap (LLM skill analysis), pipeline (orchestration), with centralized configuration and logging
- Implemented production-grade error handling with graceful degradation — try/except on every LLM call, None-check propagation, partial pipeline results, and file-based structured logging with severity levels
- Built multi-document RAG system using LangChain for generating personalized career intelligence (interview prep, study plans, readiness scoring) from retrieved job context
- Developed AI agent capabilities using LangChain Agents with tool-use patterns for autonomous career analysis workflows
- Created REST API layer using FastAPI and containerized the full application with Docker for production deployment

PROFESSIONAL EXPERIENCE

Data Analyst & BI Developer

Mar 2025 – Present

Fipsar Solutions, Chennai

- Architected and independently delivered a first-of-its-kind write-back reporting dashboard with real-time data capture using Qlik Sense Automation and Snowflake — designed the complete end-to-end architecture from research to production deployment
- Built and deployed a Global Field Medical Leadership dashboard serving stakeholders across US and EMEA regions — owned the entire data lifecycle including extraction, transformation, data modeling, and interactive visualization with KPIs, filters, and drill-downs
- Designed and delivered multiple production-ready enterprise BI applications for international clients — Channel Management dashboard (shipping, inventory, multi-domain data integration) and Quarterly Review dashboard (leadership KPIs, performance trends)
- Wrote optimized SQL queries for data validation, model accuracy checks, and logic debugging across large-scale datasets; performed SIT/UAT to ensure business alignment and production readiness
- Delivered all applications on time with zero critical defects; provided post-production hypercare support; collaborated with business stakeholders to translate feedback into functional and visual enhancements

EDUCATION

B.Tech in Computer Science

2024 | CGPA: 7.68/10

Mahaguru Institute of Technology, Kerala

SELF-DIRECTED TRAINING & CERTIFICATION

- **100-Day Generative AI Engineering Program** (Self-Directed, Project-Based) — Structured curriculum covering: LLM internals and Transformer architecture → Prompt engineering and structured output → Embeddings and vector databases → RAG systems with LangChain → AI agents and tool use → API development with FastAPI → Docker containerization. All concepts applied through hands-on project (JobLens).